

# Erlang NIFs: Native Implemented Functions

Derek Brown

January 20, 2015

# Materials

- Code examples and slides:  
[https://github.com/derek121/nif\\_demo](https://github.com/derek121/nif_demo)

# Topics

- Benefits and dangers
- The minimal, “Hello, Joe” NIF
- Handling basic data structures
- Memory management
- Scheduler issues
- Dirty schedulers

# Benefits

- Introduced in R13B03 (November 2009)
- Execute C code in-process with the VM
- Speed
- Easy way to use existing C libraries

# Dangers/Issues

- Runs in-process with the VM
- Requires knowledge of/coding in C

# A Minimal NIF

Erlang:

- Call to `erlang:load_nif/2`
- Stub (or alternate) function in the Erlang module

# A Minimal NIF

Erlang:

- Call to `erlang:load_nif/2`
- Stub (or alternate) function in the Erlang module

C:

- Declare the NIF functions
- Implementation with proper signature

# Erlang - Load

```
-on_load(init/0).
```

```
init() ->
```

```
    ok = erlang:load_nif("priv/demo_nif", 0).
```



# Erlang - Stub

```
-on_load(init/0).
```

```
init() ->
```

```
    ok = erlang:load_nif("priv/demo_nif", 0).
```

```
hello() ->
```

```
    erlang:nif_error(nif_not_loaded).
```

```
%% Or an alternate implementation in
```

```
%% Erlang
```

# C - Init

```
#include "erl_nif.h"
```

```
ERL_NIF_INIT(demo, nif_funcs,  
             load, reload, upgrade, unload)
```

load, reload (*deprecated*), upgrade, unload:

Pointers to functions called on those events. May be NULL.

# C - Init

```
#include "erl_nif.h"
```

```
ERL_NIF_INIT(demo, nif_funcs,  
             load, reload, upgrade, unload)
```

```
static ErlNifFunc nif_funcs[] = {  
    {"hello", 0, hello_nif}  
}
```

# C - Implementation

```
static ERL_NIF_TERM hello_nif(  
    ErlNifEnv* env,  
    int argc, const ERL_NIF_TERM argv[])  
{  
    return enif_make_tuple2(  
        env,  
        enif_make_atom(env, "ok"),  
        enif_make_atom(env, "joe"));  
}
```

# Build and Run

```
$ rebar compile
```

- Compiles `c_src/ .c` to `priv/`
- `port_specs` in `rebar.config` to define source-to-`.so` mappings
- `port_env` allows specifying compile/link flags

```
$ rebar shell
```

- Adds `ebin/` and `deps/*/ebin/` to codepath

# Demo Example

[https://github.com/derek121/nif\\_demo](https://github.com/derek121/nif_demo)

- “minimal” example

# Returning Values – Ints, Atoms, Strings

```
ERL_NIF_TERM t;  
t = enif_make_int(env, 24);  
  
t = enif_make_atom(env, "hello");  
  
t = enif_make_string(env, "hello",  
    ERL_NIF_LATIN1);
```

# Returning Values – Tuples

```
ERL_NIF_TERM enif_make_tuple(ErlNifEnv* env, unsigned cnt, ...)
```

```
ERL_NIF_TERM enif_make_tuple1(ErlNifEnv* env, ERL_NIF_TERM e1)
```

```
ERL_NIF_TERM enif_make_tuple2(ErlNifEnv* env, ERL_NIF_TERM e1,  
ERL_NIF_TERM e2)
```

```
ERL_NIF_TERM enif_make_tuple3(ErlNifEnv* env, ERL_NIF_TERM e1,  
ERL_NIF_TERM e2, ERL_NIF_TERM e3)
```

```
ERL_NIF_TERM enif_make_tuple4(ErlNifEnv* env, ERL_NIF_TERM e1, ...,  
ERL_NIF_TERM e4)
```

```
ERL_NIF_TERM enif_make_tuple5(ErlNifEnv* env, ERL_NIF_TERM e1, ...,  
ERL_NIF_TERM e5)
```

```
ERL_NIF_TERM enif_make_tuple6(ErlNifEnv* env, ERL_NIF_TERM e1, ...,  
ERL_NIF_TERM e6)
```

```
ERL_NIF_TERM enif_make_tuple7(ErlNifEnv* env, ERL_NIF_TERM e1, ...,  
ERL_NIF_TERM e7)
```

```
ERL_NIF_TERM enif_make_tuple8(ErlNifEnv* env, ERL_NIF_TERM e1, ...,  
ERL_NIF_TERM e8)
```

```
ERL_NIF_TERM enif_make_tuple9(ErlNifEnv* env, ERL_NIF_TERM e1, ...,  
ERL_NIF_TERM e9)
```



# Returning Values - Lists

```
ERL_NIF_TERM i1 = enif_make_int(env, 26);  
ERL_NIF_TERM i2 = enif_make_int(env, 36);  
ERL_NIF_TERM i3 = enif_make_int(env, 46);  
  
ERL_NIF_TERM list = enif_make_list3(env,  
i1, i2, i3);
```

# Receiving Values- Numeric

```
int val;  
enif_get_int(env, argv[0], &val);
```

# Receiving Values- Atoms

```
unsigned int len;  
enif_get_atom_length(env, argv[0], &len,  
    ERL_NIF_LATIN1))  
char *buf = (char *)enif_alloc(len + 1);
```

# Receiving Values- Atoms

```
unsigned int len;  
enif_get_atom_length(env, argv[0], &len,  
    ERL_NIF_LATIN1))  
char *buf = (char *)enif_alloc(len + 1);  
  
enif_get_atom(env, argv[0], buf, len + 1,  
    ERL_NIF_LATIN1);  
printf("%s\r\n", buf);  
  
enif_free(buf);
```

# Receiving Values- Strings

```
unsigned int len;  
enif_get_list_length(env, argv[0], &len);  
char *buf = (char *)enif_alloc(len + 1);
```

# Receiving Values- Strings

```
unsigned int len;  
enif_get_list_length(env, argv[0], &len);  
char *buf = (char *)enif_alloc(len + 1);  
  
enif_get_string(env, argv[0], buf, len + 1,  
    ERL_NIF_LATIN1);  
printf("%s\r\n", buf);  
  
enif_free(buf);
```

# Receiving Values- Tuples

```
const ERL_NIF_TERM *tuple;  
int arity;  
enif_get_tuple(env, argv[0], &arity,  
    &tuple);
```

# Receiving Values- Tuples

```
const ERL_NIF_TERM *tuple;  
int arity;  
enif_get_tuple(env, argv[0], &arity,  
    &tuple);  
  
for (int i = 0; i < arity; i++) {  
    int val;  
    enif_get_int(env, tuple[i], &val);  
    printf("%d: %d\r\n", i, val);  
}
```



# Receiving Values- Lists

```
ERL_NIF_TERM head, tail;  
int val;  
ERL_NIF_TERM list = argv[0];  
  
while (enif_get_list_cell(env, list, &head,  
    &tail)) {  
    enif_get_int(env, head, &val);  
    printf("%d\r\n", val);  
    list = tail;  
}
```

# Receiving Values- Lists

```
ERL_NIF_TERM head, tail;  
int val;  
ERL_NIF_TERM list = argv[0];  
  
while (enif_get_list_cell(env, list, &head,  
    &tail)) {  
    enif_get_int(env, head, &val);  
    printf("%d\r\n", val);  
    list = tail;  
}
```

# Receiving Values - Lists

```
ERL_NIF_TERM head, tail;  
int val;  
ERL_NIF_TERM list = argv[0];  
  
while (enif_get_list_cell(env, list, &head,  
    &tail)) {  
    enif_get_int(env, head, &val);  
    printf("%d\r\n", val);  
    list = tail;  
}
```

# Memory Management

- Manually create and destroy
  - E.g., `enif_alloc()`
- Resources
  - Released automatically

# Resources

- An opaque pointer passed between C and Erlang
- Deallocated after call to release and all referencing terms have been garbage collected

# Resources

- Specified via `init()`

```
ERL_NIF_INIT(demo, nif_funcs,  
             load, NULL, NULL, NULL)
```

```
int (*load)(  
    ErlNifEnv* env,  
    void** priv_data,  
    ERL_NIF_TERM load_info)
```

# Resources - Definition

```
static int load(ErlNifEnv* env, void** priv_data,
ERL_NIF_TERM load_info)
{
    ErlNifResourceType *resource_type =
        enif_open_resource_type(
            env,
            NULL, /* module_str: must be NULL */
            "resource",
            resource_dtor,
            (ErlNifResourceFlags)(ERL_NIF_RT_CREATE),
            NULL /* Used on failure */
        );

    *priv_data = resource_type;
    return 0;
}
```

# Resources - Allocating

- `ERL_NIF_TERM enif_make_resource(ErlNifEnv*env, void* obj)`
  - Allocates
- `void enif_release_resource(void* obj)`
  - If called before return, resource will be deallocated once unreferenced on Erlang side
  - Else, may be kept on NIF side to be released later



# Resources - Allocating

```
static ERL_NIF_TERM return_resource_nif(  
    ErlNifEnv* env, int argc, const ERL_NIF_TERM argv[])  
{  
    ErlNifResourceType* resource_type =  
        (ErlNifResourceType*)enif_priv_data(env);  
  
    Vals *vals =  
        (Vals *)enif_alloc_resource(  
            resource_type, sizeof(Vals));  
    vals->val1 = 23;  
    vals->val2 = 25;  
  
    ERL_NIF_TERM result = enif_make_resource(env, vals);  
    enif_release_resource(vals);  
    return result;  
}
```

# Resources - Destructors

- Called when an instance of given resource type is released
- Allows manual cleanup of allocated memory, etc.

```
typedef void ErlNifResourceDtor(  
    ErlNifEnv* env, void* obj);
```

# Example

- erlsha2

<https://github.com/vinoski/erlsha2.git>

# Scheduling

- Scheduler per core
- Processes are given a fixed number of reductions before being preempted (2000)
  - E.g., function calls, sending messages, etc.
- NIF operations are not included
- NIFs therefore need to complete within  $\sim 1$  ms

# NIF Scheduling

- `enif_schedule_nif/6`
  - Manually call with smaller tasks
- `enif_consume_timeslice/2`
  - Give scheduler hints on progress
- Example: `bitwise:exor_yield/2`

# Dirty Schedulers

- Separated to avoid interfering with regular process scheduling
- Two types: CPU-bound and IO-bound
  - Number of CPU-bound dirty schedulers  $\leq$  number of normal schedulers
  - Number of IO-bound dirty schedulers not limited
- Set via command line flags to `erl` and via `erlang:system_flag/2`

# Run Queues

- Normal schedulers (one per core)
  - Separate run queues for each
- CPU-bound schedulers
  - Share a run queue
- IO-bound schedulers
  - Share a run queue

# NIF Dirty Scheduling – Dynamic

- Pass flags to `enif_schedule_nif/6`
  - `ERL_NIF_DIRTY_JOB_CPU_BOUND`
  - `ERL_NIF_DIRTY_JOB_IO_BOUND`
- Allows decision on dirtiness or not to be made per-call



# NIF Dirty Scheduling – Static

- Specify in the ErlNifFunc array entry
  - ERL\_NIF\_DIRTY\_JOB\_CPU\_BOUND
  - ERL\_NIF\_DIRTY\_JOB\_IO\_BOUND
- E.g.:

```
static ErlNifFunc funcs[] = {  
    {"exor_dirty", 2, exor,  
     ERL_NIF_DIRTY_JOB_CPU_BOUND}};
```

# Long Scheduling

- Receive notification when regular scheduler thread runs past threshold

```
1> erlang:system_monitor(self(),
    [{long_schedule, 10}]).
{<0.38.0>, [{long_schedule, 10}]}
2> spawn(fun() -> demo:delay(1) end).
<0.66.0>
3> flush().
Shell got {monitor,<0.66.0>,long_schedule,
    [{timeout,1001},
    {in,undefined},{out,undefined}]}
ok
```

# Example

- bitwise

<https://github.com/vinoski/bitwise>

# References

API:

<http://www.erlang.org/doc/tutorial/nif.html>

[http://www.erlang.org/doc/man/erl\\_nif.html](http://www.erlang.org/doc/man/erl_nif.html)

erlsha2:

<https://github.com/vinoski/erlsha2>

bitwise:

<https://github.com/vinoski/bitwise>

Steve Vinoski 2014 Codemesh talk on schedulers

<http://www.codemesh.io/codemesh2014/steve-vinoski>

# Thanks

Thanks to Steve Vinoski for details on scheduling and dirty schedulers, and for the creation of the erlsha2 and bitwise projects.

# Questions?